

IDAWI a parallel/distributed library tailored to Research in IOT/fog/edge/distributed computing

Luc Hogie

Cnrs/Inria/Université Côte d'Azur

January 27, 2021

What is this all about?

This is about IDAWI a new library for parallel/distributed computing. IDAWI s being designed to:

- meet the needs of our on-going/future R&D projects
 - investigate IOT/decentralized algorithms (Map/Liquori)
 - provide scientific services to our applications (Gemoc)
 - compute very large graphs (Coati)
- besides, be helpful to other labs (Open Source license)
- avoid the pitfalls of existing tools

<https://github.com/lhogie/idawi>

Latest product of a software suite?

IDAWI enefits from the experience we gained in past devs:

BigGrph set of modules dedicated to distributed graph computing. Supported by Inria. Now discontinued.

JMR disk-based distributed batch processor.

JMaxGraph multi-thread graph library for processing large graphs.

MPI4Lectures multi-thread message-passing library dedicated to teaching at University of Luxembourg

Also, IDAWI earns a lot from ProActive, a distributed implementation of the Fractal specification developed by another team in our lab.

In a few words...

IDAWI enables its user to execute software components on networked computers. These components have message-oriented communication facilities that embed built-in and user-defined services will interact with each other by sending/receiving messages. Also they expose . Doing this, they distributed application made of services within the components.

Components: what they *are*

There is no global consensus on what a component is. In IDAWI a component is an object augmented with the following features:

- it organizes, along with other components, in a multi-graph
- it is able send/receive *messages* to/from other components, using a variety of transport protocols
- it forwards incoming messages
- it exposes *services* which extend its functionality

Components: what they *are used for*

IDAWI components can be used to several purposes:

- exposing computational resources in a cluster
- simulating complex systems by representing domain-specific entities and their interoperations
- providing services in a distributed application

Built-in services

Ranging from simplest to more sophisticated:

- `exit` shuts down the entire distributed system

- `ping pong` emulates the ubiquitous ping command

- `error log` archives and disseminates internal errors

- `service lifecycle` starts/stops services remotely

- `routing` guides messages traveling across the graph

Built-in services

Ranging from simplest to more sophisticated:

bencher provides performance information about the host computer

maps constructs a graph representing the topology of the network

deployment enables new components to be started

time series database stores and serves numerical time-based information on user-defined metrics

publish-subscribe notifies subscriber components of new publications on particular topics

Service

A service exposes to other components functionality for a particular concern.

- it proposes a high-level API for communication
- it defines a set of *operations*

operations are what perform user-defined computations

- has a name to executable code
- triggered by *messages* which carries parameters
- by default, operations are all executed in *parallel*
- at runtime, they can send messages (temporary results, current state, progress information, etc)

Just like components and services, operations can be created/deleted at runtime.

Message

- carries a content (that can be anything)
- to a set of components defines in a "to:" address
- has an optional "reply-to:" (that is automatically fed in the case of a synchronous emission)

An address consists of:

- an optional set of recipients components (unicast if $|R| = 1$, multicast if $|R| > 1$, broadcast if $R = null$)
- a recipient service that is expected on all recipients components
- a recipient operation

Communication

- message emission is asynchronous
- no guaranty of delivery
- if a message has "reply-to" information, the sender obtains a queue that will store response messages.

Queues enable:

- asynchronous on-the-fly processing of incoming messages
- synchronous collect and classification

Communication protocols

IDAWI currently support the following transport protocols:

- TCP: ensures reception
- UDP: is quick
- IPC: connect components in child/parent processes
 - locally, mostly useful for tests
 - remotely through SSH, across NATs and firewalls
- intra-process: method calls enables large local simulations

Deployment of components

New components can be deployed anywhere a SSH connection is possible (and rsync is available). When you

- ① determines shared file systems among computers
- ② update binaries, which includes:
 - ① incremental update of Java bytecode (using rsync)
 - ② installs the right JVM if necessary
- ③ executes a JVM and starts a component in it

The parent component initially communicate with the new component through the standard I/O streams of the SSH local process. If the new component is declared to be autonomous, it remains alive even if the I/O streams get closed.

Interoperable with external tools using REST

- the REST service launches a HTTP server
- which serve JSON documents provided by specific REST operations in services
- the REST interface gives access to the component system as a whole, regardless of which component exposes it
- executing a REST operation is done via the URL:

`http://host:port/component/service/operation/parms1,param2,...,param3`

Conclusion: most notable features

- ① multihop mobile overlay network of components
- ② deployment of new remote/local component
- ③ synchronous and asynchronous communications
- ④ support for unicast/multicast/broadcast
- ⑤ reactive message programming and stream processing
- ⑥ massively parallel computations
- ⑦ multi-protocol, REST/JSON interface
- ⑧ many base services for platform management and demo

Adding/deleting operations

operations can be declared by:

calling `Service.registerNewOperation()` allows adding/deleting new operations are runtime

adding annotated methods improves readability by isolating the code, allowing the specification of a return and parameters type. Such an operation can be unregistered but its implementation code remains, and can still be invoked from within its class.

Sending a message

operations can be declared by:

```
To to = new To(  
    MyResultType r = (MyResultType) send(content ,  
                                           Set.of(myRemoteComponent) ,  
                                           myService , myoperation )  
        .setTimeout(5)  
        .collect()  
        .ensureResults(1)  
        .first().content ;
```

Conclusion